

Applying Learning Classifier Systems to Acoustic Scene Classification: DCASE 2017 Challenge

CITS4404 Artificial Intelligence & Adaptive Systems Team Project

Yiyang Gao (21263128), Aaron Hurst (21325887), Kevin Kuek (21307006), and
Scott McCormack (21875529)

School of Computer Science and Software Engineering

3rd November, 2017

Abstract

This will be our abstract

$$\frac{\text{deletionThreshold}}{3.05} \left(\pi/2 - \arctan \left(\frac{\text{age} - \text{ageThreshold}}{20} \right) \right)$$

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

1 Introduction

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Cu-

rabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

2 Background

This section provides a brief review of learning classifier systems (2.1), the DCASE Challenge (2.2) and acoustic scene classification (2.3).

2.1 Learning Classifier Systems

First introduced in the mid-1970s, Learning Classifier Systems (LCSs) are a rule-based machine learning algorithm with a unique combination of learning mechanisms [1]. The core of

LCSs is the population of rules, or *classifiers*, which collectively form the solution to the given problem [2]. This population of classifiers is gradually evolved toward an optimal and *optimally general* set [2].

The motivation for this structure is that, when modelling and attempting to predict the outcome of complex systems, a desirable approach is to develop a distributed population of classifiers – or rules – that together form an accurate model [2, p. 2]. Each classifier, then, spans a subspace of the problem, with the population spanning the entire problem. Individual classifiers consist of a condition-action rule which says: 'If a problem instance matches this condition, perform this action'. Classifier fitness is evaluated based on feedback from the problem (also known as the environment).

The learning process of a LCS includes a rule discovery method known as covering, a generalisation-pressure effect called subsumption, a generic algorithm (GA), fitness-based deletion to maintain a finite-sized population and, in some cases, reinforcement learning [3].

The rule discovery method, covering, is used to initialise the population by adding a new classifier whenever a problem instance matches no existing classifier. Subsumption is used to eliminate classifiers with more specific conditions that are no more accurate than others with equivalent, but more general conditions. The GA is used as a secondary rule discovery method that only operates on high performing rules. Deletion is employed to maintain a finite population size by removing poorly performing classifiers. Reinforcement learning may be used as a final step in the learning cycle for problems where feedback from the environment is delayed.

A seminal work in the field of LCSs was the introduction of the eXtended Classifier System by Wilson [4, 5]. This incorporated a number of features which substantially improved the performance of the LCS [4].

A key distinction amongst LCS applications is between supervised learning, in which feedback from the environment is delayed (such as robot navigation), and offline learning, where

feedback is immediate and the correct action known in advance (such as classification tasks). This distinction determines whether reinforcement learning is necessary and affects how classifier accuracy is calculated.

2.2 DCASE Challenge

Sound classification, or machine listening, is seen as a promising research field with multifarious applications [6]. The DCASE challenge has been established to encourage and support work in this space by providing standardised

what (sound recognition, machine listening), why (motivations) and how (mechanics of challenge)

baseline solution, best results

2.3 Acoustic Scene Classification

more detailed description of this particular task from the challenge

3 Experiments

blah

3.1 Modification of Existing Code

Is there a better title you'd suggest Scott?

Description of Ryan's code, how we modified it and the results obtained

3.2 Adapted XCS(R) Design

Details of the LCS we made, specific design changes relative to the standard implementation for our problem

4 Results

4.1 Environment Representations

Alternative feature processing investigated

4.2 Parameter Tuning

Reason why standard deviation and mean are used, what information they carry

5 Discussion

Car and office confusion matrix Plot for std dev and mean (for car and office) to explain Car, office and city centre (with confusion and plots) Then add metro station and show results (worse) - explain why it is worse Show a classifier on a graph to show how it can match instances Overall accuracy (all 15 features)

6 Conclusion

References

- [1] M. V. Butz, “Learning classifier systems,” in *Springer Handbook of Computational Intelligence*. Springer, 2015, pp. 961–981.
- [2] R. J. Urbanowicz and J. H. Moore, “Learning Classifier Systems: A Complete Introduction, Review, and Roadmap,” *Journal of Artificial Evolution and Applications*, vol. 2009, pp. 1–25, 2009.
- [3] M. V. Butz and S. W. Wilson, “An Algorithmic Description of XCS,” in *International Workshop on Learning Classifier Systems*. Springer, 2000, pp. 253–272.
- [4] P. L. Lanzi, “Learning Classifier Systems: Then and Now,” *Evolutionary Intelligence*, vol. 1, no. 1, pp. 63–82, 2008.
- [5] O. Sigaud and S. W. Wilson, “Learning Classifier Systems: A Survey,” *Soft Computing*, vol. 11, no. 11, pp. 1065–1078, 2007.
- [6] Mesaros, Annamaria and Heittola, Toni and Dimint, Aleksandr and Elizalde, Benjamin and Shah, Ankit and Vincent, Emmanuel and Raj, Bhiksha and Virtanen, Tuomas, “DCASE 2017 challenge setup: Tasks, datasets and baseline system,” in *Detection and Classification of Acoustic Scenes and Events*, 2017.